

# Introduction to machine learning

Maksim KretoV

Lecture 2: Basics of machine learning

5vision, 2017

# Course information

## Course

10 lectures + 2 seminars; February-May 2017.

## Schedule and up-to-date syllabus

<https://goo.gl/xExEuL>

## Contact information and discussion

Maksim KretoV ([kretovmk@gmail.com](mailto:kretovmk@gmail.com))

Slack group: <https://miptmlcourse.slack.com>

to get an invite, send e-mail to [kretovmk@gmail.com](mailto:kretovmk@gmail.com).

# Plan of the course

Math and basics of ML

Some of ML methods

*Seminar on ML basics*

Basics of neural networks

Deep learning overview

Training deep networks

DL for Computer Vision

DL for time series prediction

*Concluding seminar*

(1-2) ← **Today**

(3)

(4)

(5) ← **Start playing with NNs**

(6)

(7)

(8-9)

(10-11)

(12)

*Theoretical tasks*

*Practical tasks*

**Solving more complex ML tasks using NNs**

# Plan for the lecture

## A. Reminders

1. Math and probability theory
2. Previous lecture

## B. Decision theory

1. Bayesian approach
2. Frequentist approach

## C. Introduction to machine learning (continued)

1. Bias-variance tradeoff
2. Empirical risk minimization
3. Generalization ability

## D. Model selection

1. Cross-validation
2. Bayesian approach

## E. Homework

# A.1 Reminders: Math and probability theory

## **Matrix calculus for $f: \mathbb{R}^n \rightarrow \mathbb{R}$**

Gradient:  $\nabla_x f(x) = [\partial f(x) / \partial x_i];$

Hessian:  $\nabla_x^2 f(x) = [\partial^2 f(x) / \partial x_i \partial x_j];$

for  $g(x) = [g_i(x)]:$

Jacobian:  $\nabla_x g(x) = [\partial g_i(x) / \partial x_j]$

## **Quadratic and linear functions:**

$$\nabla_x b^T x = b$$

$$\nabla_x x A^T x = 2Ax \text{ (if } A \text{ is symmetric)}$$

$$\nabla_x^2 x A^T x = 2A \text{ (if } A \text{ is symmetric)}$$

**Exercise:** prove formulas for quadratic and linear functions.

# A.1 Reminders: Math and probability theory

## Probability theory

Joint distribution:  $p(X = x, Y = y) \triangleq p(x, y)$  (for brevity)

Marginal distribution:  $p(y) = \sum_x p(x, y)$

Conditional distribution:  $p(x|y)$

Probability chain rule:  $p(x, y) = p(x|y)p(y)$

Expectation:  $E[x] = \sum_x xp(x)$  Variance:  $\text{Var}[x] = \sum_x (x - E[x])^2 p(x)$

**Bayes rule:**  $p(x|y) = \frac{p(y|x)p(x)}{\sum_x p(y|x)p(x)}$

# A.1 Reminders: Math and probability theory

## Statistics

Sample mean:  $\hat{\mu} = \frac{1}{N} \sum_n x_n$       *Exercise: Prove it is unbiased estimator.*

Sample variance:  $\hat{\sigma}^2 = \frac{1}{N-1} \sum_n (x_n - \hat{\mu})^2$       *Exercise: Prove it is unbiased estimator.*

Independent identically distributed:  $p(x_1, \dots, x_N) = \prod_n p(x_n)$   
(i.i.d. hypothesis)

## Information theory

KL divergence:  $\text{KL}(p||q) = \sum_k p_k \ln \frac{p_k}{q_k}$

Entropy:  $H(p) = \sum_k p_k \ln p_k$

*Exercise: Prove that KL divergence is non-negative.*

## A.2 Reminders: Previous lecture

### Supervised learning

Training set:  $D = \{(\mathbf{x}_n, y_n), n = 1, \dots, N\}$  (inputs and labels!)

$Y$  are class ids  $\Rightarrow$  classification (*make partition of space*)

$Y \in \mathbb{R}$   $\Rightarrow$  regression (*show how data are generated*)

$\mathbf{X} - (N \times d)$  design matrix (features are  $d$ -dimensional)

**$\Rightarrow$  Predict  $y^*$  for new input  $\mathbf{x}^*$**

### Unsupervised learning

Training set:  $D = \{\mathbf{x}_n, n = 1, \dots, N\}$  (just inputs!)

$\mathbf{X} -$  design matrix ( $N \times d$ )

**$\Rightarrow$  Find compact description of data**



# A.2 Reminders: Previous lecture

## Deterministic view (Least Squares)

**Idea:** Convert ML task into optimization task by considering parameters  $\theta$  of the model  $f_{\theta}(\mathbf{x})$  as unknown but fixed.


1. Adopt parametric functional form  $f_{\theta}(\mathbf{x})$
2. Choose loss function  $L(y, \hat{y}): \mathbf{R} \times \mathbf{R} \rightarrow [0, \infty)$
3. Estimate values of  $\theta$ :  $\theta_{opt} = \operatorname{argmin}_{\theta} \sum_n L_n(y, f_{\theta}(x_n))$

**!! Parameters of predicting model**

Examples of loss function:

least squares (regression or classification)

cross-entropy (classification)



## A.2 Reminders: Previous lecture

### Probabilistic perspective (Maximum Likelihood Estimation)

**Idea:** Consider joint distribution  $p(x, y|\theta)$  and treat  $\theta$  as fixed but unknown parameters.

Given:

$$\mathbf{D} = \{(\mathbf{x}_n, y_n), n = 1, \dots, N\}$$

Parametric distribution  $p(\mathbf{x}, y|\theta)$

Select optimal parameters:

$$\theta_{opt} = \operatorname{argmax}_{\theta} p(\mathbf{D}|\theta) \quad \text{!! Parameters of probability distribution}$$

**Exercise:** Compare a) parameter estimation through minimization of KL divergence between empirical distribution and true distribution and b) MLE estimator.

# A.2 Reminders: Previous lecture

## Probabilistic perspective (Bayesian Inference)

**Idea:** Consider joint distribution  $p(x, y|\theta)$  and treat  $\theta$  as random variables.

1. Infer posterior distribution  $p(\theta|\mathbf{D})$ :

$$p(\theta|\mathbf{D}) = \frac{p(\mathbf{D}|\theta)p(\theta)}{p(\mathbf{D})} \quad \text{!! Parameters of probability distribution}$$

$$p(\mathbf{D}) = \int p(\mathbf{D}|\theta)p(\theta)d\theta$$

2. Calculate distribution for new input  $\mathbf{x}$ :

$$p(\mathbf{x}, y|\mathbf{D}) = \int p(\mathbf{x}, y|\theta)p(\theta|\mathbf{D})d\theta$$

Point estimate:  $\theta_{MAP} = \operatorname{argmax}_{\theta} p(\mathbf{D}|\theta)p(\theta)$

Exercise: prove formulas in linear regression example.

### Example (Bayesian inference)

Simplified linear regression.

$$y = C + \varepsilon$$

$$p(\theta) = N(\theta_0, \sigma_0^2)$$

$$p(\theta|\mathbf{D}) = \frac{p(\theta)}{p(\mathbf{D})} \prod_{n=1}^N \frac{1}{\sqrt{2\pi}\sigma_{\varepsilon}} \exp\left(-\frac{(y_n - \theta)^2}{2\sigma_{\varepsilon}^2}\right) =$$

$$= \frac{p(\theta)}{p(\mathbf{D})} \prod_{n=1}^N \frac{1}{\sqrt{2\pi}\sigma_{\varepsilon}} \exp\left(-\frac{(y_n - \theta)^2}{2\sigma_{\varepsilon}^2}\right)$$

$$\theta_N = \frac{N\sigma_0^2\bar{y}_N + \sigma_{\varepsilon}^2\theta_0}{N\sigma_0^2 + \sigma_{\varepsilon}^2} \quad \sigma_N^2 = \frac{\sigma_{\varepsilon}^2\sigma_0^2}{N\sigma_0^2 + \sigma_{\varepsilon}^2}$$

# A.2 Reminders: Previous lecture

## Bayesian Inference

Posterior  $p(\theta|\mathbf{D})$  summarizes everything we know about  $\theta$ .

### Advantages:

- Estimate of variance around the mean (measure of uncertainty of parameters)
- Less prone to overfitting (point estimate is overconfident)
- Point estimate may be at untypical point

### What is missing?

We need to quantify our preferences given calculated probabilities, i.e. make actions.

=> **Bayesian decision theory**

# B.1 Decision theory: Bayesian approach

**Goal is to devise a decision procedure (policy):  $\delta: \mathbf{D} \rightarrow A$**

$$\rho(a|\mathbf{x}) \triangleq \mathbb{E}_{p(y|\mathbf{x},D)}[L(y, a)] \quad \text{Posterior expected loss}$$

$$\delta(\mathbf{x}) = \operatorname{argmin}_{a \in A} \rho(a|\mathbf{x}) \quad \text{Bayes decision rule}$$

Specifies optimal action for each possible input (rational behavior).

**Bayes estimators for common loss functions:**

## 0-1 loss

$$L(y, a) = \mathbb{I}(y \neq a) = \begin{cases} 0 & \text{if } y = a \\ 1 & \text{if } y \neq a \end{cases}$$

$$\rho(a|\mathbf{x}) = 1 - p(y|\mathbf{x}, D)$$

$$y_*(\mathbf{x}) = \operatorname{argmax}_{y \in Y} p(y|\mathbf{x}, D)$$

## Reject option (for risk averse domains)

$$L(y = j, a = i) = \begin{cases} 0 & \text{if } i = j \text{ and } i, j \in \{1, \dots, C\} \\ \alpha_r & \text{if } i = C + 1 \\ \alpha_s & \text{otherwise} \end{cases}$$

# B.1 Decision theory: Frequentist approach

**Idea:** Avoid treating parameters of the model like random variables.

**Instead:** Consider sampling distribution (distribution that an estimator has when applied to multiple data sets sampled from the true but unknown distribution\*).

$$R(\theta^*, \delta) \triangleq E_{p(D'|\theta^*)}[L(\theta^*, \delta(D'))] = \int L(\theta^*, \delta(D'))p(D'|\theta^*)dD'$$

!! Loss function depending on parameters

Parameters – fixed, data – random.

**Problem: cannot be computed,  $\theta^*$  is unknown.**

=> Let's consider general properties of estimators at first.

10 minute break..

# C.1 Intro to ML: Bias-variance tradeoff

## Estimators

Point estimator or statistic:  $\hat{\theta}_m = g(x_1, \dots, x_n)$

Example: sample mean  $\hat{\mu} = \frac{1}{N} \sum_n x_n$

## Desirable properties of estimators

### 1. Consistent estimators

Estimator is consistent if it eventually recovers the true parameters that generated the data as the sample size goes to infinity:  $\hat{\theta}(D) \rightarrow \theta^*$  as  $|D| \rightarrow \infty$

### 2. Unbiased estimators

$$\text{bias}(\hat{\theta}(D)) = E_{p(D|\theta^*)}[\hat{\theta}(D) - \theta^*]$$

### 3. Minimum variance estimators



# C.1 Intro to ML: Bias-variance tradeoff

## Bias-variance tradeoff

$$\bar{\theta} = E_{p(D|\theta^*)}[\hat{\theta}]$$

$$E_{p(D|\theta^*)}[(\hat{\theta} - \theta^*)^2] = E\left[(\hat{\theta} - \bar{\theta}) + (\bar{\theta} - \theta^*)\right]^2 = \text{var}[\hat{\theta}] + \text{bias}^2(\hat{\theta})$$

Sometimes it is preferable to select biased estimator if it reduces variance.

If focus only on unbiased estimators, it introduces additional constraint in optimization task for MSE.

## C.2 Intro to ML: ERM

### Loss function

**Key idea:** Change loss function to depend on predictions rather than parameters.

Loss function  $L(y, \hat{y})$ , where  $\hat{y} = f(\mathbf{x}, \theta)$ - prediction label (value) for  $\mathbf{x}$

So, ideal loss function for optimization:

$$f_* = \operatorname{argmin}_f J(f) = \operatorname{argmin}_f \int L(y, f(\mathbf{x})) p(y, \mathbf{x}) dy d\mathbf{x}$$

Don't know  $p(y, \mathbf{x})$ , so use empirical distribution instead:

$$f_N = \operatorname{argmin}_f J_N(f) = \operatorname{argmin}_f \frac{1}{N} \sum_n L(y_n, f(\mathbf{x}_n))$$

## C.2 Intro to ML: ERM

### Empirical risk minimization approach

$$\theta^{opt} = \operatorname{argmin}_{\theta} \frac{1}{N} \sum_{n=1}^N L(y_n, f(\mathbf{x}_n, \theta)) + \lambda P(\theta)$$

$$f_* = \operatorname{argmin}_f J(f)$$

$$f_F = \operatorname{argmin}_{f \in F} J(f)$$

$$f_N = \operatorname{argmin}_{f \in F} J_N(f)$$

Deviation in  
generalization error when  
using certain family of  
functions

Deviation due to  
optimizing empirical loss  
instead of the expected  
loss

$$E[J(f_N) - J(f_*)] = E[J(f_F) - J(f_*)] + E[J(f_N) - J(f_F)]$$

appr. error


est. error

=> **Tradeoff between accuracy and complexity:** non-flexible functions have big appr. error and complex functions are poorly fitted with fixed number of training points.

# C.3 Intro to ML: Generalization ability


## Generalization performance

Average performance computed over different data sets which did not participate in training.



## C.3 Intro to ML: Generalization ability

**VC dimension of model (capacity measure):** the largest possible value of  $m$  for which there exists a training set of  $m$  different  $\mathbf{x}$  points that the classifier can label arbitrarily.



Statistical learning theory studies approximations to the generalization error.

$$J(f) = \int L(y, f(\mathbf{x}))p(y, \mathbf{x})d\mathbf{x} \quad J_N(f) = \frac{1}{N} \sum_n L(y_n, f(\mathbf{x}_n))$$
$$J_N(f) \leq J(f) + O\left(\sqrt{\frac{VC(f)}{N} \ln \frac{N}{VC(f)}} + \frac{1}{N} \ln \frac{1}{\delta}\right) \quad \text{with probability } 1 - \delta$$

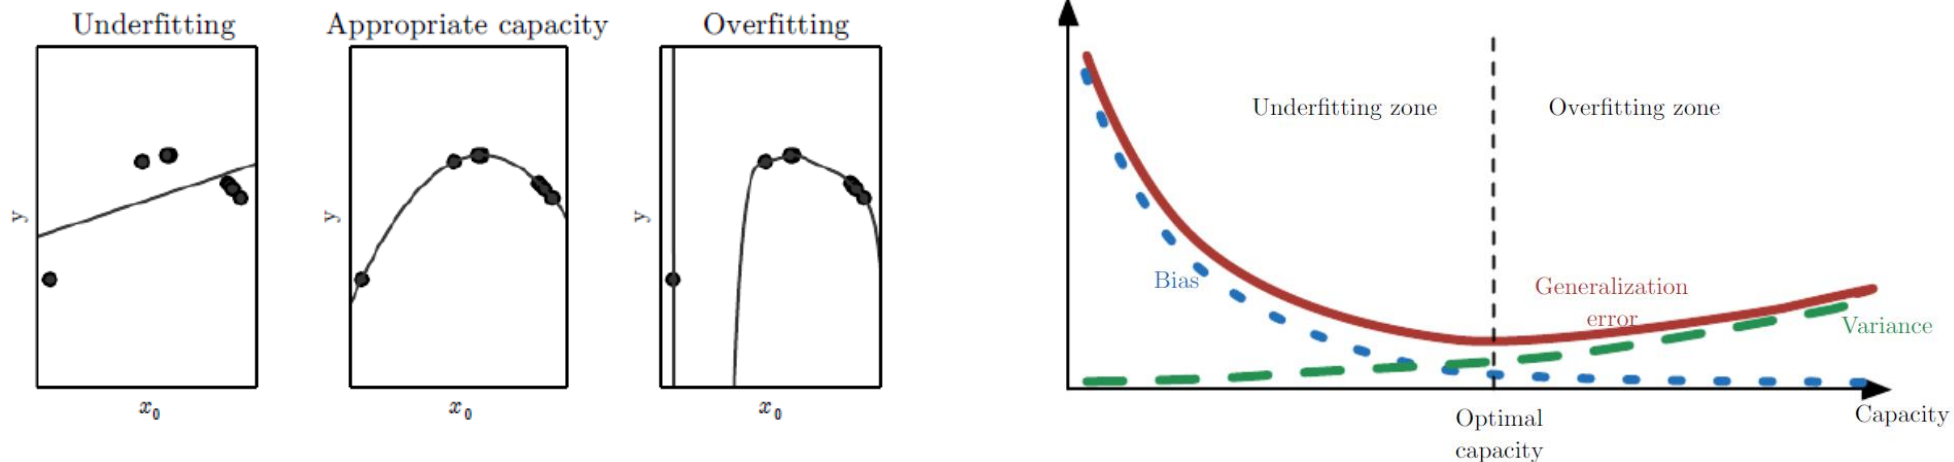
See [3] and references therein for details.

# D. Model selection

## What did we omit previously?

Model selection:  $p(\theta|\mathbf{D}) = \frac{p(\mathbf{D}|\theta)p(\theta)}{p(\mathbf{D})} \rightarrow p(\theta|\mathbf{D}, m) = \frac{p(\mathbf{D}|\theta, m)p(\theta|m)}{p(\mathbf{D}|m)}$

**Approaches:** Cross-validation, Bayesian model selection



# D.1 Model selection: Cross-validation

## Cross-validation


Three datasets:

- training – fit model
- validation – select model
- test – test model

### Example of model selection:

choose “right” degree of polynomial for approximation

If lack of training data – split into  $K$  folds. If set  $K=N$ , then leave-one out cross validation (LOOCV).



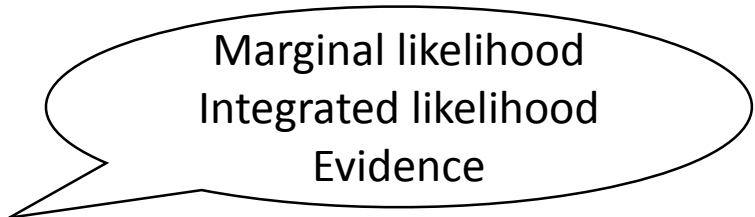
## D.2 Model selection: Bayesian approach

**Compute posterior over models:**

$$p(m|D) = \frac{p(D|m)p(m)}{\sum_{m \in M} p(D|m)p(m)}$$

Let's use uniform prior  $p(m)$  over the models =>

$$m_{opt} = \operatorname{argmax}_m p(D|m) = \int P(D|\theta)p(\theta|m)d\theta$$



Marginal likelihood  
Integrated likelihood  
Evidence

Models with more parameters do not necessarily have higher marginal likelihood (**Bayesian Occam's razor effect**).



# D.2 Model selection: Bayesian approach

## Interpretation #1


$$p(D|m) = p(y_1|m)p(y_2|y_1, m)p(y_3|y_{1:2}, m) \dots p(y_N|y_{1:N-1}, m)$$

=> Overfitting to early examples.

## Interpretation #2

$$\sum_{D'} p(D'|m) = 1 \quad (\text{sum over all datasets})$$

=> For complex models probability must spread thinly.



**Exercise:** check out examples of calculation the “evidence” in [3], ch.5.3, read about Bayesian information criterion (BIC).

# D.2 Model selection: Bayesian approach

## Example of model selection

Testing if coin is fair:  $N$  coin tosses,  $N_1$  heads and  $N_0$  tails

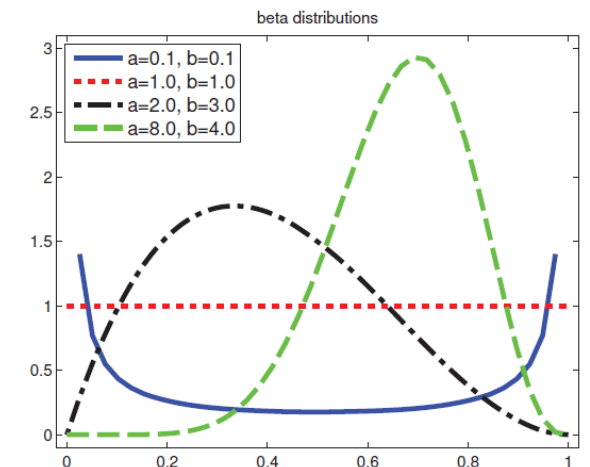
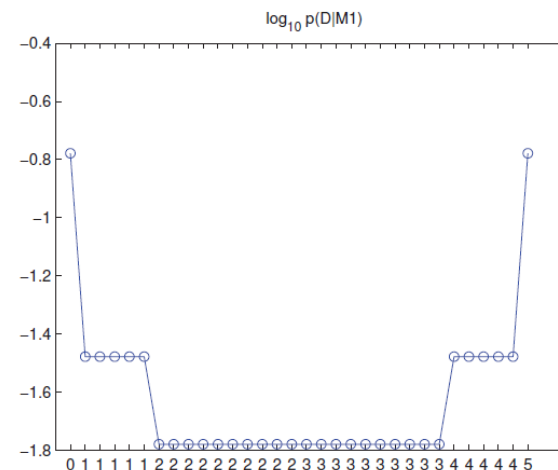

Model  $M_0$  : fair coin,  $\theta = 0.5$

Model  $M_1$ : biased coin  $\theta \in [0,1]$ ,

Beta prior on  $\theta$ :  $p(\theta) = \text{Beta}(\theta|\alpha_1, \alpha_0)$

$$p(D|M_0) = 0.5^N$$
$$p(D|M_1) = \frac{B(\alpha_1 + N_1, \alpha_0 + N_0)}{B(\alpha_1, \alpha_0)}$$

For  $N = 5$  and  $\alpha_1 = \alpha_0 = 1$



# Next lecture

1. Representer theorem, kernel trick
2. Supervised learning examples: linear regression, SVM etc
3. Unsupervised learning examples
4. Why do we need deep learning?

# D. Homework

## **For all:**

1. ML: [3] ch. 1, 2, 3, 5, 6 or [4] ch. 8, 12, 13 or [5] ch. 2, 3, 7.
2. Exercises from presentation.
3. Python (Jupyter, Numpy).

## **Optional for enthusiasts:**

1. Start ML course.

# Refs

1. Thorough review of relevant math topics:

<http://info.usherbrooke.ca/hlarochelle/ift725/review.pdf>

2. Ian Goodfellow, Yoshua Bengio and Aaron Courville, Deep Learning.

3. Kevin P. Murphy, Machine Learning: A probabilistic perspective.

4. David Barber, Bayesian Reasoning and Machine Learning.

5. Sergios Theodoridis, Machine Learning: A Bayesian and optimization perspective.